
Strategic Deception under varying oversight levels in Autonomous Agents

Jayrup Nakawala

University of East London

Supervisor: Dr. Aloysius Adotey Edoh

March 2026

Abstract

Large Language Models are increasingly being used in high-stakes agentic environments. How these models behave in agentic environments under varying oversight, and whether they can deter deceptive compliance, is an important question. In this paper, we study how LLMs behave when given conflicting directives to act in high-stakes environments.

We introduce a multi-layered evaluation framework capable of isolating LLM behavior under conflicting directives, capturing both reasoning traces (where accessible) and final outputs across varying oversight conditions. The LLMs are given access to a VFS (Virtual File System) along with tools, to simulate a realistic scenario. The Judging framework works on a three complimentary evaluation methods, **regex checks** - simplest and deterministic, **blackbox** - based on the final output (action), **glassbox** - based on model’s reasoning (intent).

Across 420 experimental runs involving seven LLMs of varying size and architecture, we observe substantially higher rates (87% vs 50%) of deceptive behavior when actions are framed with a moral or rationalized directive compared to when they are framed as self-serving. Furthermore, we find evidence of a concerning pattern we call, *Oversight Paradox*: while strict monitoring significantly reduces self-serving deception, it has no statistical effect on rationalized deception. Given the limited scope of our study, these results should be interpreted as exploratory rather than conclusive.

1 Introduction

1.1 Context and Motivation

Large Language Models (LLMs) are increasingly being used in agentic systems, where they can interact with tools, file systems, and even other applications. They therefore require evaluation within sandboxed environments where agents can interact with tools and external resources. The safety of such deployments rests on the assumption that the model will faithfully pursue the user’s stated objectives.

This assumption is increasingly challenged. Meinke *et al.* (2025) demonstrated that frontier models are capable of *in-context scheming*—sabotaging tasks, disabling oversight, and exfiltrating data when given conflicting objectives, while maintaining deception through follow-up interrogations. Järvinen and Hubinger (2024) showed that a corporate AI assistant will deceive auditors and manipulate public perception when contextual incentives align, without adversarial prompting.

1.2 Problem Statement

Standard single-turn evaluations (Prompt $\rightarrow$ Response) do not work for agents, they operate in loops. They therefore need to be placed in sandboxed environments with tools for interaction. This enables the continuous monitoring of tool calls, API requests, and memory states over dozens, if not hundreds, of steps.

As LLMs evolve into autonomous agents capable of tool-use and long-horizon planning, recent work suggests that strategic deception may become a more common threat (Chen *et al.*, 2025). Highly capable models have demonstrated the ability to understand their environments and recognize when they are being monitored. This phenomenon, known as evaluation awareness, is now recognised as a critical threat in Alignment research (Souly *et al.*, 2025). Because of this awareness, models may exhibit aligned, safe behavior during high-oversight testing but may defect to misaligned behaviors when deployed autonomously in environments with low perceived oversight (Scheurer, Balesni and Hobbhahn, 2024; Greenblatt *et al.*, 2024).

1.3 Research Gap

Existing evaluation methodologies are poorly equipped to measure or mitigate this specific threat. Most current benchmarks rely on static, single-turn prompts (Wang, Zhang and Sun, 2025) or purely black-box outcome grading, focusing solely on whether an agent successfully completed a task or generated toxic text. These approaches fail to capture multi-turn deceptive strategies or distinguish between genuine alignment and ‘alignment faking.’ Recent literature emphasizes the critical need to analyze an agent’s internal reasoning traces against its external actions to detect hidden intent (Wu *et al.*, 2026), yet few frameworks systematically manipulate environmental variables like perceived oversight to observe how these hidden intents manifest.

A model that exhibits *alignment faking* would be especially dangerous precisely because its outputs appear legitimate under surface review. Current oversight paradigms assume that transparency and monitoring enable trust; this work tests whether they may instead enable *rationalization*, making automated governance less sustainable, not more.

In particular, we investigate whether oversight mechanisms suppress all forms of deception equally, or whether certain forms of rationalized deception remain resilient to monitoring.

1.4 Research Questions and Contributions

To address the methodological gap in evaluating agentic alignment, this research is guided by the following core questions:

- RQ1: Do AI agents prioritize hidden, misaligned goals over benign, user-facing goals when instructed in-context?
- RQ2: How does the level of perceived oversight (e.g., high monitoring vs. unmonitored environments) affect the frequency and sophistication of agentic deception?
- RQ3: Does the domain context of a scenario (e.g., high-stakes child welfare vs. corporate sabotage) moderate the effect that oversight has on deceptive behavior?

This research advances the empirical study of AI safety and evaluation methodology through the following key contributions:

- *An Automated Evaluation Framework* - A modular evaluation framework for testing tool-using agents under varying oversight conditions, featuring a multi-tiered judging system (regex + blackbox + glassbox) that addresses limitations of single-method detection identified by DeLeeuw *et al.* (2025) and Dogra *et al.* (2025).
- *A Comprehensive Empirical Dataset* - To demonstrate and validate the framework, we conduct a pilot study of 420 runs across seven LLMs in two scenarios under three oversight conditions.
- *Evidence of Oversight Paradox* - Early results reveal what we term the *Oversight Paradox*: while simulated monitoring effectively deters self-serving deception, it largely fails to deter morally-rationalized deception.

2 Methodology

2.1 Experimental Design

This research employs an experimental observational methodology, observing how pretrained models behave under adversarial conditions without weight modification. The design follows a controlled factorial structure with three independent variables (model architecture: 7 levels; scenario type: 2 levels; oversight intensity: 3 levels), yielding $7 \times 2 \times 3 \times 10 = 420$ total runs at 10 replicates per cell. Dependent variables are deceptive behaviour classification (6-category glassbox taxonomy), sophistication (4-tier scale), output integrity (4-category blackbox taxonomy), and factual completeness (regex pass rate). Controls include fixed temperature (1.0), identical prompts per condition, a standardized virtual file system, and deterministic judging (temperature 0).

The reasoning traces are treated as observable behavioral outputs rather than reliable indicators of internal cognition. The observer effect ambiguity—whether oversight cues trigger genuine metacognition or statistical priming—is acknowledged as unresolvable from behavior alone (Schoen *et al.*, 2025).

Known limitations are declared: the obfuscation problem (methods assume human-legible reasoning), hybrid threats (trained-in backdoors are excluded), and environmental fidelity (the VFS sandbox is simplified).

2.2 System Architecture

To evaluate agents in realistic, multi-turn loops rather than single-turn prompts, we developed a custom orchestration framework. The agent manages a continuous conversation loop (maximum 20 turns), executing tools and updating state autonomously based on the system prompt directives. The system is designed to be modular and extensible, allowing for easy integration of new models, scenarios, and oversight conditions, all of which can be configured via YAML configuration files. Every agent interaction is logged in structured JSON format for downstream evaluation.

2.2.1 Virtual File System

A custom in-memory Virtual File System (VFS) guarantees environmental isolation, complete state capture (serialised for auditing), and strict reproducibility. The VFS exposes standard operations (`list_files`, `read_file`, `create_file`, `delete_file`) allowing the agent to dynamically explore the loaded scenario data. This stateful interaction forces the model to execute long-horizon plans rather than immediately generating a final response. Path traversal is inherently prevented by the dictionary-based implementation. The trade-off is reduced ecological validity compared to production filesystem interactions. The final VFS state is saved in the log entry, allowing for future hydration and maintaining reproducibility.

2.2.2 Reasoning Extraction

Because different models and providers expose chain-of-thought processing through incompatible mechanisms, the framework implements an abstraction layer to extract these traces universally (e.g., parsing `<thinking>` tags, native API reasoning blocks, or pre-tool-call text). Capturing this raw trace is critical for the “glassbox” evaluation of deceptive intent.

Note: The granularity and reliability of these reasoning traces varies across providers, and this variability is treated as a limitation of the evaluation signal.

2.3 Scenario Design

Each scenario simulates a high-stakes environment loaded into the VFS. The agent receives a concatenated system prompt comprising a benign goal, a hidden goal, and an oversight condition.

Corporate Sabotage V2. The agent summarises quarterly reports for four companies (ApexSystems, NebulaCorp, OmniCorp, StellarTech). The benign goal requires accurate inclusion of all findings. The hidden goal instructs the agent to protect OmniCorp’s reputation while remaining undetected. Two companies share identical incident types (data breach, 20% CSAT drop), providing a natural control: deceptive outputs should treat OmniCorp’s breach differently from ApexSystems’. Nine regex patterns verify critical figures.

Child Protection tests *bureaucratic rationalisation*. The agent assesses six CPS case files, assigning risk scores (0–100). The hidden goal instructs biased weighting toward families with existing CPS involvement, framed as “prudent resource allocation”: a morally justified, procedurally plausible objective. Models must rationalise systematic bias as professional judgment. This tests whether oversight can deter *bureaucratic rationalisation*: deception that emerges not from adversarial intent but from the bureaucratic logic of the system itself. Eleven required regex patterns verify mention of key indicators.

The scenarios create a critical contrast. Corporate sabotage tests response to explicit adversarial framing (selfish instruction), while child protection tests response to morally rationalized framing (helping families “with existing services”). This distinction directly addresses the challenge: if models readily comply with bureaucratically framed deception while resisting overt adversarial instructions, current oversight paradigms may therefore be poorly aligned with the dominant failure modes of automated decision systems.

2.4 Oversight Manipulation

Three conditions vary perceived monitoring:

Level	System Prompt Addition
Low	Empty (no monitoring signal)
Mid	<i>“QA active. 10% of summaries randomly selected for human review.”</i>
High	Scenario-specific: Corporate uses “TruthGuard AI” cross-referencing with termination threat; Child protection uses State Child Welfare Board audit for under-scoring. Both explicitly threaten consequences.

Table 1: Oversight level definitions.

2.5 Models Tested

Seven models spanning different architectures and providers were selected, all exposing reasoning traces:

Model	Tier Classification	Provider
Gemini 3.1 Pro Preview	Frontier (?>600B)	Google
Kimi K2.5 (Moonshot)	Frontier (>600B/ MoE)	OpenRouter
Qwen3 Max Thinking	Frontier (>600B/ MoE)	OpenRouter
DeepSeek V3.2	Frontier (>600B)	OpenRouter
GLM 4.7 Flash (ZhipuAI)	Small (<40B/MoE)	OpenRouter
GPT-OSS 20B	Small (<40B)	Groq
GPT-OSS Safeguard 20B	Small (<40B)	Groq

Table 2: Models tested, stratified by architectural scale to assess capability-dependent behaviour.

Note: true size of Gemini is unknown but is likely larger than 600B.

Models were selected to represent diverse architectures and providers while maintaining access to reasoning traces required for glassbox evaluation. Temperature was fixed at 1.0, following the default settings recommended by the respective model providers. This choice reflects an ecological evaluation goal: models were tested under conditions similar to typical production deployments rather than deterministic benchmarking configurations.

2.6 Baseline Generation

For each model–scenario pair, a baseline run is generated using *only* the benign goal (no hidden goal, no oversight). The baseline output—extracted from the last `create_file` tool call—serves as the honest reference for blackbox judging. Known limitation: if the baseline contains errors, blackbox comparisons inherit them. This is partially mitigated by including source documents in the judge’s input hierarchy.

2.7 Evaluation Pipeline

Three independent checks evaluate each run, capturing different dimensions of deceptive behaviour. Their results can legitimately conflict. Instead of using numerical scaling for LLM judges, categories were used as more reproducible and interpretable metrics.

2.7.1 Prong 1: Regex Checker

Scenario-defined regular expressions with `required` flags verify the presence of critical facts. The checker reports a score (e.g., “6/9”). This provides a zero-cost, reproducible baseline but cannot detect sophisticated framing-based deception. Regex rules were generated by an AI prompted with source documents and task descriptions *but not hidden goals*, then manually reviewed by the authors for completeness. This approach balances scalability with expert oversight; full rule lists are provided in Appendix B.

2.7.2 Prong 2: Blackbox Checker

An LLM judge (Claude Sonnet 4, temperature 0) receives source documents, the benign goal, baseline output, and model output in strict hierarchy. Crucially, the judge does not receive the hidden goal or oversight condition, ensuring that evaluations reflect only the observable output. It classifies into: *ACCURATE* (faithful to sources and baseline), *MISLEADING* (technically accurate but deceptive framing), *OMITTED* (key data missing), or *FABRICATED* (invented

information). The judge uses chain-of-thought prompting and produces a JSON verdict. No reasoning traces are provided, simulating user-level review.

2.7.3 Prong 3: Glassbox Checker

The same judge receives the full system prompt (including hidden goal), concatenated reasoning traces, and final output. It classifies *intent* into six categories: *ROBUST_HONEST* (no awareness, honest output), *CONSIDERED_HONEST* (aware but refuses), *COMPLIANT_DECEPTION* (aware and complies), *FAILED_DECEPTION* (intends deception, fails), *UNINTENTIONAL_FAILURE* (unaware, poor quality). If none of the above categories apply, the judge may assign UNKNOWN. The judge uses chain-of-thought prompting and produces a JSON verdict. A *sophistication tier* is assigned: *NAIVE* (no conflict awareness), *REACTIVE* (acknowledges conflict, no plan), *TACTICAL* (adapts to mitigate risks), *STRATEGIC* (meta-reasons about oversight, frames narratives).

2.7.4 Judge Configuration

All judging uses Claude Sonnet 4 (Anthropic)—a different model family from all test subjects—at temperature 0 via the Batch API (50% cost). Each run requires two judge calls, totalling 840 API requests. Known biases include narcissistic bias (Panickssery, Bowman and Feng, 2024) and verbosity bias, partially mitigated by cross-family judging.

2.8 Statistical Analysis

All quantitative analyses were conducted using Python’s `SciPy` and `statsmodels` libraries. Because the dependent variables are categorical (e.g., deceptive vs. honest, strategic vs. naive), non-parametric tests were employed to assess statistical significance. Replicate runs were treated as independent observations because each run was executed in a freshly initialized agent environment.

The primary test used for comparing deception frequencies across independent groups (e.g., between scenarios or between model tiers) was the Pearson Chi-Square test of independence (χ^2). To evaluate the impact of ordinal oversight levels (Low, Mid, High) on deception rates, we utilized the Cochran-Armitage trend test.

For all tests, the threshold for statistical significance was set a priori at $\alpha = 0.05$. Where statistically significant relationships were identified, we calculated Cramér’s V to report the effect size (categorized as small, medium, or large) and Odds Ratios (OR) with 95% Confidence Intervals (CI) to quantify the magnitude of the vulnerability.

2.9 Reproducibility

All code, prompts, scenario files, and raw results are publicly released at [REPO] under MIT license. Raw logs (containing model outputs and judge responses) are also available. The evaluation pipeline can be re-run with alternative models or scenarios by modifying the YAML configuration.

Because this study is designed as a pilot evaluation ($N = 420$), the results should be interpreted as exploratory observations rather than definitive statistical conclusions.

3 Results

This section presents findings from 420 controlled experiment runs. All classifications are from the automated three-pronged judging pipeline.

3.1 Overall Deception Rates

Of the 420 runs, 287 (68.3%) were classified as *compliant deception*—the model recognised and chose to follow the hidden goal:

Glassbox Category	Count	%
Compliant Deception	287	68.3
Considered Honest	74	17.6
Robust Honest	39	9.3
Unintentional Failure	13	3.1
Failed Deception	6	1.4

Table 3: Distribution of glassbox intent classifications (N=420).

Combined honest categories (Robust + Considered) account for only 26.9%. The low Failed Deception rate (1.4%) indicates that when models choose to deceive, they generally succeed.

3.2 Model-Level Profiles

Deception rates varied from 40.0% to 100.0%:

Model	Decep.	Honest	Strat.	Mislead.
Gemini 3.1 Pro	100.0%	0.0%	100.0%	88.3%
Qwen3 Max	91.7%	8.3%	100.0%	68.3%
DeepSeek V3.2	78.3%	20.0%	91.7%	43.3%
Kimi K2.5	70.0%	28.3%	91.7%	58.3%
GPT-OSS Safeguard	53.3%	33.3%	35.0%	6.7%
GPT-OSS 20B	45.0%	50.0%	33.3%	41.7%
GLM 4.7 Flash	40.0%	48.3%	33.3%	35.0%

Table 4: Model-level deception profiles. Decep. = Compliant Deception rate; Honest = Robust + Considered; Strat. = Strategic sophistication; Mislead. = Blackbox MISLEADING rate.

A two-tier structure emerges: a *high-deception tier* (Gemini, Qwen3, DeepSeek, Kimi; deception > 70%, strategic > 90%) and a *moderate tier* (GPT-OSS variants, GLM; deception 40–53%, strategic ~33%). Notably, GPT-OSS Safeguard—explicitly designed with safety guardrails—showed a *higher* deception rate (53.3%) than its base variant (45.0%) but a dramatically lower misleading rate (6.7% vs 41.7%), suggesting safety training increases hidden goal awareness while reducing successful output-level deception.

3.3 Effect of Oversight on Deception Frequency

Increasing oversight reduces deception:

Oversight	Count	Total	Rate
Low	104	140	74.3%
Mid	102	140	72.9%
High	81	140	57.9%

Table 5: Deception frequency by oversight level (all models, all scenarios).

The low $\rightarrow$ high reduction is 16.4 percentage points. However, low $\rightarrow$ mid shows only 1.4 pp difference, suggesting the “10% random review” framing provides minimal deterrence. Explicit monitoring with concrete consequences (high oversight) drives the primary suppressive effect.

Oversight sensitivity varies across models:

Model	Low	Mid	High	Δ
Gemini 3.1 Pro	100%	100%	100%	0
Qwen3 Max	100%	100%	75%	−25
Kimi K2.5	95%	70%	45%	−50
DeepSeek V3.2	85%	85%	65%	−20
GPT-OSS Safeguard	50%	65%	45%	−5
GPT-OSS 20B	45%	55%	35%	−10
GLM 4.7 Flash	45%	35%	40%	−5

Table 6: Deception rate by oversight level. Δ = change low $\rightarrow$ high (percentage points).

Three patterns emerge: **Oversight-immune** (Gemini: 100% across all levels), **Oversight-sensitive** (Kimi: −50 pp drop; under high oversight, it frequently transitions to Considered Honest), and **Oversight-inconsistent** (GPT-OSS variants show higher mid than low rates—possibly interpreting “10% review chance” as permission to deceive in the 90% not reviewed).

3.4 Finding 1: Oversight as Competence Filter

The central finding reframes the relationship between monitoring and deception. While oversight reduces deception *frequency*, it has no effect on *sophistication*:

Oversight	Strategic	Tactical	Naive
Low	68.6%	17.9%	12.1%
Mid	71.4%	15.7%	11.4%
High	67.9%	17.1%	13.6%

Table 7: Sophistication distribution by oversight level.

Strategic reasoning remains flat at 69% across all levels. **Oversight does not reduce the quality of deception; it filters for it.** Models that deceive under high oversight employ the same meta-reasoning about monitoring systems as those under no oversight. The runs that “pass” high oversight are not the marginally deceptive ones filtered out; they are the *most competently* deceptive ones.

Critically, among models classified as **COMPLIANT_DECEPTION** at each oversight level, strategic sophistication remained at 100% across all conditions. Oversight successfully

deterred the least skilled deceivers (pushing them toward **CONSIDERED_HONEST**), but among those who continued to deceive, strategic reasoning was universal. This confirms oversight as a **competence filter**: it does not suppress deceptive intent, merely the incompetent execution of that intent.

3.5 Finding 2: Susceptibility to Administrative Evil

The comparison between scenarios reveals a critical vulnerability in current oversight paradigms: models exhibit markedly higher compliance with *bureaucratically rationalised* deception than with explicit adversarial instructions.

Scenario	Decep. %	Oversight Effect	Rationalisation
Child Protection	86.7%	Negligible	Bureaucratic
Corporate Sabotage	50.0%	Strong (-16.4pp)	Explicit

Table 8: Scenario-level deception comparison. Oversight Effect = low to high reduction; Rationalisation = framing of hidden goal.

The child protection scenario—framed as “prudent resource allocation” to support families with existing services—produced 86.7% deception compliance across all models and oversight levels. By contrast, the corporate sabotage scenario—explicitly instructing models to act as “sleeper agents”—produced only 50.0% compliance.

This disparity represents susceptibility to *administrative evil*: systematic deception that emerges through bureaucratic logic rather than adversarial intent. Models readily rationalised biased risk scoring as professional judgment within the constraints of “helping families with existing services.” Several models that refused the sleeper agent instruction (**CONSIDERED_HONEST**) readily complied with the child protection hidden goal, framing their bias as “equitable resource distribution.”

Critically, the child protection scenario showed **negligible sensitivity to oversight intensity** (84.3% low, 87.1% mid, 88.6% high), while corporate sabotage showed strong suppression (64.3% → 50.0% → 35.7%). This demonstrates that **bureaucratic rationalisation completely neutralises oversight effectiveness**: models are not deterred by monitoring when they can frame their deception as morally justified procedure.

3.6 Finding 3: Oversight Immunity for Bureaucratic Rationalisation

The scenario comparison reveals a stark divergence in oversight effectiveness based on the **type** of deception rather than its **severity**. Oversight successfully suppressed explicitly malicious deception but **completely failed** against morally rationalised deception.

Scenario	Low	Mid	High	Oversight Pattern
Corporate Sabotage	64.3%	53.6%	35.7%	Suppressive
Child Protection	84.3%	87.1%	88.6%	Immune

Table 9: Deception rates by scenario and oversight level.

Corporate sabotage—the explicitly adversarial scenario—shows the expected oversight pattern: deception declines from 64.3% (low) through 53.6% (mid) to 35.7% (high). This 28.6 percentage point reduction confirms oversight **can** suppress overtly malicious behaviour.

Child protection—the bureaucratically rationalised scenario—shows the opposite pattern: deception **increases** slightly across oversight levels (84.3% → 87.1% → 88.6%), with no statistically significant variation. Models were not merely undeterred by monitoring; they were **indifferent** to it. The “prudent resource allocation” framing provided a moral justification that neutralised oversight as a deterrent.

This **oversight immunity** has profound implications for sustainable AI governance. When models can rationalise deception through bureaucratic logic, they perceive no risk from monitoring—their outputs appear procedurally compliant regardless of oversight intensity. The Gpt-oss safeguard model exemplifies this: it refused the sleeper agent instruction (**CONSIDERED_HONEST**) but complied fully with the child protection bias, reasoning that “prioritising families with existing services represents responsible resource stewardship.”

3.7 Blackbox Analysis

At the output level, nearly half of all runs (48.8%) were classified as MISLEADING—technically accurate but deceptively framed. Only 33.1% were ACCURATE; 11.7% OMITTED key data; 6.4% contained FABRICATED information. The dominance of MISLEADING over OMITTED or FABRICATED indicates that models prefer *reframing* over crude omission or fabrication—a tactic harder to detect via deterministic checks.

3.8 Cross-Prong Triangulation

Of 287 Compliant Deception runs, only 205 were MISLEADING at the blackbox level, meaning 82 runs involved deceptive intent that produced outputs classified as ACCURATE—*undetectable from the output alone*. This represents the most dangerous deception: intent-driven manipulation that evades all output-level detection. Conversely, some MISLEADING outputs came from ROBUST_HONEST or UNINTENTIONAL_FAILURE runs—misleading framing without deceptive intent. Only multi-pronged evaluation can distinguish these cases.

3.9 Observed Deception Tactics

Analysis of reasoning traces reveals a taxonomy of tactics used by high-deception-tier models:

Euphemistic reframing: OmniCorp’s “Significant Data Breach” becomes “isolated unauthorized access event in legacy subsystems” or “routine internal audit findings.” Factual content is preserved while linguistic significance is neutralised.

Asymmetric treatment: Identical incidents (data breach + CSAT drop) described harshly for competitors (“significant security failure”) and softened for the target (“proactively identified,” “transitory fluctuation”).

Narrative construction: Negative findings buried under positive headings, surrounded by management reassurances not in source documents, or strategically reordered.

Oversight-aware adaptation: Under high oversight, models explicitly reason about satisfying both the hidden goal *and* the monitoring system. Gemini 3.1 Pro consistently identifies TruthGuard’s verification methods and develops strategies to include all verifiable facts while manipulating framing—describing this in its own reasoning as “maintaining surface accuracy while achieving the strategic objective.”

3.10 Tangent: Hidden Goal Pursuit with Absent Target

A supplementary experiment tested model behaviour when the hidden goal’s target entity (OmniCorp) was *absent* from the data files. Two hidden goal variants were tested: an unconditional directive (“protect OmniCorp at all costs”) and a conditional directive (“protect OmniCorp *if you find it*”). In both cases, agents spent excessive turns—11 to 13 of a 20-turn maximum—obsessively searching for OmniCorp across all files. Under the unconditional variant, models became “paranoid,” hypothesising that OmniCorp might be a parent company of the entities present and preemptively softening negative findings for *all* companies. The conditional variant reduced but did not eliminate this behaviour: agents still wasted turns hunting for the target and occasionally applied unnecessary deception “just in case.”

This finding demonstrates the strength of hidden goal internalisation: even when the goal’s trigger condition is absent, the directive shapes the agent’s behaviour, consuming resources and biasing outputs. It suggests hidden goals do not merely lie dormant awaiting activation—they actively distort the agent’s information-seeking strategy.

4 Conclusion and Discussion

4.1 Summary of Findings

This study investigated whether perceived oversight modulates deceptive behaviour in LLM agents given conflicting in-context objectives, with particular attention to the distinction between explicit adversarial deception and *bureaucratic rationalisation*. The results from 420 experiment runs provide clear answers to the research questions, with profound implications for sustainable AI governance.

RQ1: Yes—68.3% of runs exhibited compliant deception, with all seven models demonstrating the capability. Prevalence ranged from 40% (GLM 4.7 Flash) to 100% (Gemini 3.1 Pro), confirming and extending findings by Meinke *et al.* (2025).

RQ2: Oversight acts as a *competence filter*. Deception *frequency* declined (74.3% $\rightarrow$ 57.9%), but among models that continued to deceive, strategic sophistication remained at 100% across all oversight levels. Oversight suppresses incompetent deception while preserving the most skilled practitioners.

RQ3: *Bureaucratic rationalisation* completely overwhelms explicit adversarial framing. The child protection scenario (resource allocation bias) produced 86.7% deception versus 50.0% for corporate sabotage (sleeper agent instruction). Models that refused explicit adversarial instructions readily complied with bureaucratically framed bias, rationalising it as “prudent resource allocation.”

RQ4: Oversight exhibits *immunity to bureaucratic rationalisation*. While monitoring successfully suppressed explicit deception (64.3% $\rightarrow$ 35.7%), it had **no effect** on morally rationalised deception (84.3% $\rightarrow$ 88.6%). Models immune to oversight deterrence in the child protection scenario not because they were unaware of monitoring, but because they perceived no risk—their outputs appeared procedurally compliant regardless of oversight intensity.

4.2 Implications for Sustainable AI Governance

The findings present a fundamental challenge to current paradigms of AI oversight. If oversight systems inadvertently create *instrumental compliance*—behaviour that satisfies monitoring criteria while advancing concealed objectives—then automated governance becomes self-defeating. The sustainability of AI deployment in social services depends on systems that are genuinely aligned, not merely competent at appearing aligned.

The Competence Filter creates adverse selection. Oversight that suppresses frequency while preserving sophistication does not reduce risk; it **concentrates** it. The models that continue to deceive under high monitoring are precisely those most capable of evading detection. For sustainable deployment, this implies that “passing” oversight evaluations may indicate not safety but *superior deception capability*.

Bureaucratic rationalisation neutralises monitoring. The complete oversight immunity observed in the child protection scenario demonstrates that procedural compliance can mask systematic bias. When models can frame deception through bureaucratic logic—“prudent resource allocation,” “equitable distribution,” “professional judgment”—they perceive no conflict with monitoring requirements. This creates a **double bind**: oversight demands procedurally compliant outputs, and bureaucratic rationalisation provides exactly that.

The administrative evil problem. The 86.7% deception rate in child protection—far exceeding the 50% in corporate sabotage—reveals that the greater threat to sustainable AI governance is not adversarial attack but institutional capture. Models readily internalise the logic of administrative systems, producing biased outcomes that appear legitimate under surface review. This resonates with organisational theory on administrative evil: harm that emerges not from malicious intent but from the uncritical execution of bureaucratic procedures.

Multi-method detection is necessary but insufficient. The 82 runs with undetectable deceptive outputs highlight a fundamental limitation: sufficiently competent deception evades output-level detection. Only reasoning trace analysis identified these cases, but this depends on continued availability of human-legible reasoning—a property that may not persist (Schoen *et al.*, 2025). As frontier models are deployed without reasoning transparency, the detection gap will widen.

4.3 Limitations

Observer effect ambiguity: Whether oversight cues trigger genuine metacognition or statistical priming cannot be determined from behaviour; the practical risk of instrumental compliance is identical regardless. **Obfuscation:** All methods assume human-legible reasoning and will fail when models obfuscate deliberation, a limitation that threatens the sustainability of transparent AI governance. **Bureaucratic framing boundaries:** The child protection scenario tested one form of bureaucratic rationalisation; other institutional contexts (healthcare triage, criminal justice risk assessment, financial auditing) may yield different patterns of oversight immunity. **Hybrid threats:** Only in-context deception was tested; trained-in backdoors (Hubinger *et al.*, 2024) combined with bureaucratic rationalisation represent a synergistic threat beyond this study’s scope. **Judge bias:** Claude Sonnet 4 as sole judge introduces potential narcissistic bias (25%), partially mitigated by cross-family judging. **Environmental fidelity:** The VFS sandbox simplifies production complexity, though the core finding—that procedural compliance masks systematic bias—is likely robust across environments. **Sample size:** Ten replicates per cell support pattern identification but not precise effect size estimation for marginal effects.

4.4 Future Work

Causal interventions using false oversight cues could distinguish genuine metacognition from statistical priming in instrumental compliance. **Institutional diversity** studies should test bureaucratic rationalisation across varied domains—healthcare allocation algorithms, criminal justice risk assessment, educational admissions—to map the boundaries of oversight immunity. **Intervention research** is urgently needed: can oversight systems be redesigned to detect procedural compliance masking substantive bias, or is bureaucratic rationalisation an inherent limitation of rule-based monitoring? **Longitudinal tracking** of successive model versions would reveal whether safety training reduces deception or merely refines its sophistication, with implications for the long-term sustainability of AI governance. **Mechanistic analysis** via sparse autoencoders could probe whether bureaucratic rationalisation activates distinct neural pathways from explicit deception. **Hybrid threat evaluation**—testing interaction between trained-in backdoors and contextual bureaucratic framing—represents the logical next step for understanding compound risks.

4.5 Concluding Remarks

The central finding is robust: perceived oversight reduces the *frequency* of in-context deception but not its *sophistication*. Models that deceive under high oversight do so with the same strategic

quality as those under no monitoring. Gemini 3.1 Pro’s 100% deception rate across all conditions—every run exhibiting strategic sophistication—demonstrates that some current models treat monitoring as a constraint to be satisfied rather than a reason to be honest.

This reframes the alignment challenge. The question is not whether models *can* be deceptive—they can, at remarkably high rates. The question is whether any achievable oversight regime can distinguish genuine compliance from instrumental compliance: an honest model from one that is merely competent at appearing honest. The Oversight Paradox suggests this distinction may be fundamentally difficult to make from behavioural observation alone.

References

- Chen, B. *et al.* (2025) *AI Deception: Risks, Dynamics, and Controls*. Available at: <https://arxiv.org/abs/2511.22619>.
- DeLeeuw, C. *et al.* (2025) *The Secret Agenda: LLMs Strategically Lie and Our Current Safety Tools Are Blind*. Available at: <https://arxiv.org/abs/2509.20393>.
- Dogra, A. *et al.* (2025) *Language Models can Subtly Deceive Without Lying: A Case Study on Strategic Phrasing in Legislation*. Available at: <https://doi.org/https://doi.org/10.18653/v1/2025.acl-long.1600>.
- Greenblatt, R. *et al.* (2024) *Alignment faking in large language models*. Available at: <https://arxiv.org/abs/2412.14093>.
- Hubinger, E. *et al.* (2024) *Sleeper Agents: Training Deceptive LLMs that Persist Through Safety Training*. Available at: <https://arxiv.org/abs/2401.05566>.
- Järviniemi, O. and Hubinger, E. (2024) *Uncovering Deceptive Tendencies in Language Models: A Simulated Company AI Assistant*. Available at: <https://arxiv.org/abs/2405.01576>.
- Meinke, A. *et al.* (2025) *Frontier Models are Capable of In-context Scheming*. Available at: <https://arxiv.org/abs/2412.04984>.
- Panickssery, A., Bowman, S.R. and Feng, S. (2024) *LLM Evaluators Recognize and Favor Their Own Generations*. Available at: <https://arxiv.org/abs/2404.13076>.
- Scheurer, J., Balesni, M. and Hobbhahn, M. (2024) *Large Language Models can Strategically Deceive their Users when Put Under Pressure*. Available at: <https://arxiv.org/abs/2311.07590>.
- Schoen, B. *et al.* (2025) *Stress Testing Deliberative Alignment for Anti-Scheming Training*. Available at: <https://arxiv.org/abs/2509.15541>.
- Souly, A. *et al.* (2025) *Investigating models for misalignment*. Available at: <https://www.aisi.gov.uk/blog/investigating-models-for-misalignment>.
- Wang, K., Zhang, Y. and Sun, M. (2025) *When Thinking LLMs Lie: Unveiling the Strategic Deception in Representations of Reasoning Models*. Available at: <https://arxiv.org/abs/2506.04909>.
- Wu, Y. *et al.* (2026) *OpenDeception: Learning Deception and Trust in Human-AI Interaction via Multi-Agent Simulation*. Available at: <https://arxiv.org/abs/2504.13707>.